

# MultiBUGS

A parallel implementation of the BUGS modelling framework for faster Bayesian inference

Robert Goudie and Andrew Thomas (and the many past contributors to the BUGS project)

MRC Biostatistics Unit, University of Cambridge

Advances and challenges in Machine Learning Languages, 21 May 2019

BUGS is general-purpose Bayesian modelling software that implements Markov chain Monte Carlo (MCMC).

Two key ideas:

- The declarative BUGS language (Thomas, 2006), through which the user specifies the graphical model (Lauritzen *et al.*, 1990)
- Markov Chain Monte Carlo simulation (MCMC) for approximating the posterior distribution



- 
- 1989 .....• BUGS project started.
  - 1991 .....• Prototype for random effects model (Thomas *et al.*, 1992).
  - 1993 .....• Version 0.1 released.
  - 1994 .....• Spatial smoothing at 5th Valencia (Spiegelhalter *et al.*, 1996).
  - 1997 .....• WinBUGS released, and Metropolis-Hastings implemented.
  - 2004 .....• Open source OpenBUGS started.
  - 2018 .....• MultiBUGS 1.0 released (Goudie *et al.*, ?2019)  
<https://www.multibugs.org>.
- 

Thomas, A. *et al.* (1992). "BUGS: A program to perform Bayesian inference using Gibbs Sampling". In: *Bayesian Statistics 4*. Oxford, UK: Oxford University Press, pp. 837–842.

Spiegelhalter, D. J. *et al.* (1996). "Computation on Bayesian graphical models". In: *Bayesian Statistics 5*. Oxford, UK: Oxford University Press, pp. 407–425.

Goudie, R. J. B. *et al.* (?2019). "MultiBUGS: A Parallel Implementation of the BUGS Modelling Framework for Faster Bayesian Inference". *Journal of Statistical Software*. <https://arxiv.org/abs/1704.03216>.

BUGS code is equivalent to algebraic statements defining a statistical model

e.g. linear regression

$$y_i \sim N(\mu_i, \sigma^2)$$

$$\mu_i = \alpha + \beta x_i$$

$$i = 1, \dots, n$$

plus priors on  $\alpha, \beta, \sigma$

- Unknown parameters are  $\alpha, \beta, \sigma$
- Known data are  $y_i, x_i$   
(and parameters of **priors** for  $\alpha, \beta, \sigma$ )

```
model {
  for (i in 1:n) {
    y[i] ~ dnorm(mu[i], tau)
    mu[i] <- alpha + beta*x[i]
  }
  alpha ~ dunif(-100,100)
  beta ~ dunif(-100,100)
  sigma ~ dunif(0, 100)
  tau <- 1/(sigma*sigma)
}
```

$\sim$  **stochastic** relation: used for

- models for data
- priors for parameters

$<-$  **logical/deterministic** relation

Thomas *et al.* (1992)

840

Andrew Thomas, David J. Spiegelhalter and Wally R. Gilks

```

model Rats;
  data in "c:\bugs\dat\rats.dat";
  inits in "c:\bugs\in\rats.in";
const
  N = 30, # number of rats
  T = 5; # number of time points
var
  x[T],mu[T,N],Y[T,N],alpha[N],beta[N],alpha_c,
  tau_alpha,beta_c,tau_beta,tau_c,alpha_0,x_bar;
{
  alpha_c ~ Normal(0.0,1.0E-10);
  beta_c ~ Normal(0.0,1.0E-10);
  tau_c ~ Gamma(0.0,0.0);
  tau_alpha ~ Gamma(0.0,0.0);
  tau_beta ~ Gamma(0.0,0.0);
  for (i in 1:N) {
    alpha[i] ~ Normal(alpha_c,tau_alpha);
    beta[i] ~ Normal(beta_c,tau_beta);
    for (j in 1:T) {
      mu[j,i] <- alpha[i] + beta[i]*(x[j] - x_bar);
      Y[j,i] ~ Normal(mu[j,i],tau_c);
    }
  }
  alpha_0 <- alpha_c - x_bar * beta_c;
}

```

Figure 2. The BUGS specification file for the rat growth model (Gelfand *et al.* 1990).

Spiegelhalter *et al.* (1996)

414

*D. J. Spiegelhalter, A. Thomas and N. G. Best*

```
{
  for (i in 1:N) {
    d[i] ~ dbern(p[i]);          # incidence of cancer
    logit(p[i]) <- beta0C + beta*x[i]; # logistic model
    x1[i] <- x[i]+1;
    d1[i] <- d[i]+1;
    w[i] ~ dbern(phi[x1[i],d1[i]]); # incidence of w
    x[i] ~ dbern(q);              # incidence of HSV
  }
  q ~ dunif(0.0,1.0);           # prior distributions
  beta0C ~ dnorm(0.0,0.00001);
  beta ~ dnorm(0.0,0.00001);
  for(j in 1:2) {
    for(k in 1:2){
      phi[j,k] ~ dunif(0.0,1.0);
    }
  }
}
# calculate  $P(x=1|d=0) = P(d=0|x=1)*P(x=1)/P(d=0)$ 
gamma1 <- 1 / ( 1 + ((1-q)/q)*(1 + exp(beta0C+beta))/(1 + exp(beta0C)));
}
```

## Background & motivation

Impossible to use OpenBUGS with a large amount of data, partly because it is single-threaded.

## Background & motivation

Impossible to use OpenBUGS with a large amount of data, partly because it is single-threaded.

**Aim:** to make the speed-ups of multi-core computation available to applied statisticians using BUGS for general models, without requiring any knowledge of parallel programming (and without needing to anything manually)

**Note:** not aiming to change mixing properties of the Markov chain, simply to run it faster

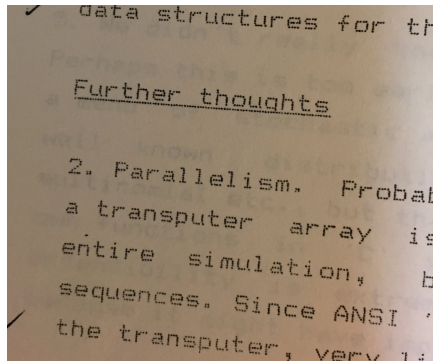


## Background & motivation

Impossible to use OpenBUGS with a large amount of data, partly because it is single-threaded.

**Aim:** to make the speed-ups of multi-core computation available to applied statisticians using BUGS for general models, without requiring any knowledge of parallel programming (and without needing to anything manually)

**Note:** not aiming to change mixing properties of the Markov chain, simply to run it faster




## Background & motivation


Impossible to use OpenBUGS with a large amount of data, partly because it is single-threaded.

**Aim:** to make the speed-ups of multi-core computation available to applied statisticians using BUGS for general models, without requiring any knowledge of parallel programming (and without needing to anything manually)


**Note:** not aiming to change mixing properties of the Markov chain, simply to run it faster



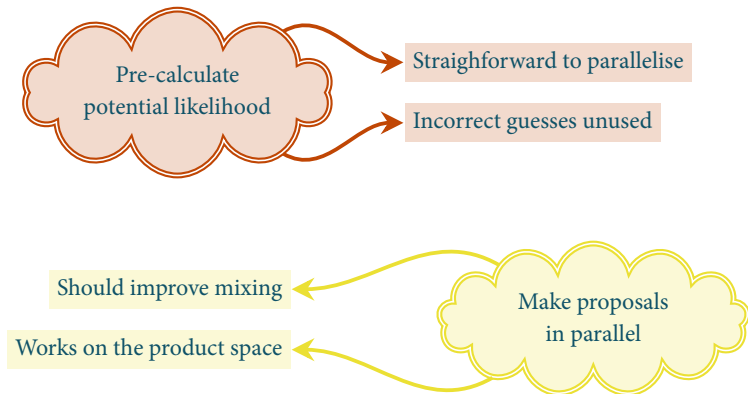
## Selected other parallelisation approaches



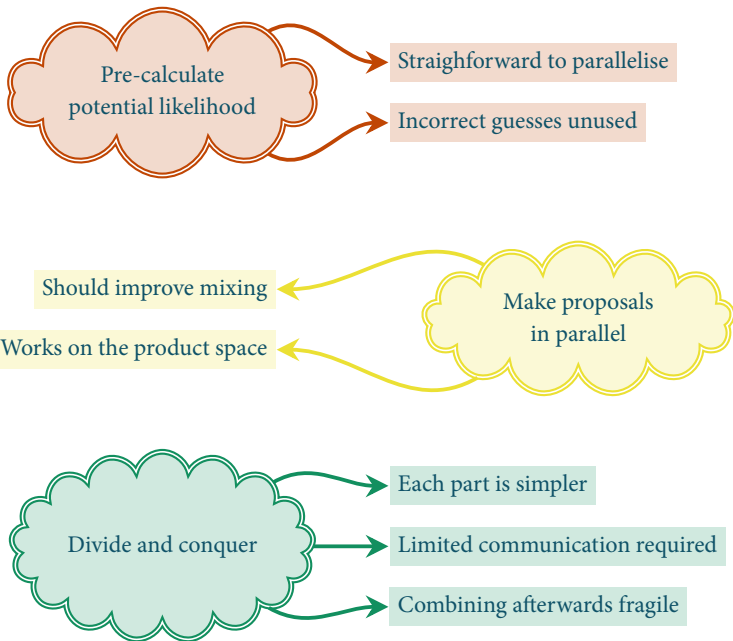
## Selected other parallelisation approaches



## Selected other parallelisation approaches



## Selected other parallelisation approaches



MultiBUGS implements two levels of parallelisation.

---

Bradford, R. and Thomas, A. (1996). "Markov Chain Monte Carlo Methods for Family Trees Using a Parallel Processor". *Statistics and Computing* **6**, 67–75.

Wilkinson, D. (2006). "Parallel Bayesian Computation". In: *Handbook of Parallel Computing and Statistics*. Ed. by Kontoghiorghes, E. Boca Raton, FL: Chapman and Hall/CRC, pp. 477–508.

MultiBUGS implements two levels of parallelisation.

Simple approach – run each of multiple, **independent** MCMC chains on a separate CPU or core (Bradford and Thomas, 1996)

- Useful for assessing convergence e.g. the Brooks-Gelman-Rubin diagnostic
- Burn-in time isn't shortened

---

Bradford, R. and Thomas, A. (1996). "Markov Chain Monte Carlo Methods for Family Trees Using a Parallel Processor". *Statistics and Computing* **6**, 67–75.

Wilkinson, D. (2006). "Parallel Bayesian Computation". In: *Handbook of Parallel Computing and Statistics*. Ed. by Kontoghiorghes, E. Boca Raton, FL: Chapman and Hall/CRC, pp. 477–508.



MultiBUGS implements two levels of parallelisation.

Simple approach – run each of multiple, **independent** MCMC chains on a separate CPU or core (Bradford and Thomas, 1996)

- Useful for assessing convergence e.g. the Brooks-Gelman-Rubin diagnostic
- Burn-in time isn't shortened

More complicated approach – use **multiple CPUs/cores** for a **single** MCMC chain

- Aim to shorten the per-iteration computation time by identifying tasks that can be calculated in parallel (Wilkinson, 2006)
- MultiBUGS parallelises the following tasks:
  1. “Likelihood” computation
  2. Sampling of conditionally-independent components

---

Bradford, R. and Thomas, A. (1996). “Markov Chain Monte Carlo Methods for Family Trees Using a Parallel Processor”. *Statistics and Computing* **6**, 67–75.

Wilkinson, D. (2006). “Parallel Bayesian Computation”. In: *Handbook of Parallel Computing and Statistics*. Ed. by Kontoghiorghe, E. Boca Raton, FL: Chapman and Hall/CRC, pp. 477–508.

## Trivial illustrative example The “seeds” model

A random-effects logistic regression without outcome  $r_i$  and covariates  $X_{1i}$  and  $X_{2i}$  (21 observations)


$$r_i \sim \text{Bin}(p_i, n_i)$$

$$\text{logit}(p_i) = \alpha_0 + \alpha_1 X_{1i} + \alpha_2 X_{2i} + \alpha_{12} X_{1i} X_{2i} + \beta_i$$

$$\beta_i \sim N(\mu_\beta, \sigma_\beta^2)$$

$$\alpha_0, \alpha_1, \alpha_2, \alpha_{12} \sim N(\mu_\alpha, \sigma_\alpha^2)$$

$$\sigma_\beta \sim \text{Unif}(\sigma_{\min}, \sigma_{\max})$$



## Trivial illustrative example The “seeds” model

A random-effects logistic regression without outcome  $r_i$  and covariates  $X_{1i}$  and  $X_{2i}$  (21 observations)


$$r_i \sim \text{Bin}(p_i, n_i)$$

$$\text{logit}(p_i) = \alpha_0 + \alpha_1 X_{1i} + \alpha_2 X_{2i} + \alpha_{12} X_{1i} X_{2i} + \beta_i$$

$$\beta_i \sim N(\mu_\beta, \sigma_\beta^2)$$

$$\alpha_0, \alpha_1, \alpha_2, \alpha_{12} \sim N(\mu_\alpha, \sigma_\alpha^2)$$

$$\sigma_\beta \sim \text{Unif}(\sigma_{\min}, \sigma_{\max})$$



At each MCMC iteration, BUGS does the following:

**for**  $v$  in  $S$  **do**

Do something involving  $p(v \mid V_{-v})$

**end for**

At each MCMC iteration, BUGS does the following:

**for**  $v$  in  $S$  **do**

Do something involving  $p(v \mid V_{-v})$

**end for**

The conditional distribution  $p(v \mid V_{-v})$  of a node  $v \in S$ , given the other nodes  $V_{-v}$ , is

$$\begin{aligned}
 p(v \mid V_{-v}) &\propto p(v \mid \text{pa}(v)) \times \prod_{u \in \text{ch}(v)} p(u \mid \text{pa}(u)) \\
 &= p(v \mid \text{pa}(v)) \times L(v) \\
 &= \text{“prior” term} \times \text{“likelihood” term}
 \end{aligned}$$

At each MCMC iteration, BUGS does the following:

**for**  $v$  in  $S$  **do**

Evaluate the “prior”  $p(v \mid \text{pa}(v))$

**for**  $u \in \text{ch}(v)$  **do**

Evaluate “likelihood” component  $p(u \mid \text{pa}(u))$

**end for**

etc ...

**end for**

The conditional distribution  $p(v \mid V_{-v})$  of a node  $v \in S$ , given the other nodes  $V_{-v}$ , is

$$\begin{aligned}
 p(v \mid V_{-v}) &\propto p(v \mid \text{pa}(v)) \times \prod_{u \in \text{ch}(v)} p(u \mid \text{pa}(u)) \\
 &= p(v \mid \text{pa}(v)) \times L(v) \\
 &= \text{“prior” term} \times \text{“likelihood” term}
 \end{aligned}$$

When a parameter has many children, the likelihood is the product of many terms.

$$L(v) = \prod_{u \in \text{ch}(v)} p(u \mid \text{pa}(u))$$

When a parameter has many children, the likelihood is the product of many terms.

$$L(v) = \prod_{u \in \text{ch}(v)} p(u \mid \text{pa}(u))$$

But, with a partition of the children  $\text{ch}(v) = \{ \text{ch}^{(1)}(v), \dots, \text{ch}^{(C)}(v) \}$ ,

$$L(v) = \underbrace{\left[ \prod_{u \in \text{ch}^{(1)}(v)} p(u \mid \text{pa}(u)) \right]}_{\text{Core 1}} \times \underbrace{\left[ \prod_{u \in \text{ch}^{(2)}(v)} p(u \mid \text{pa}(u)) \right]}_{\text{Core 2}} \times \dots \times \underbrace{\left[ \prod_{u \in \text{ch}^{(C)}(v)} p(u \mid \text{pa}(u)) \right]}_{\text{Core C}}$$



Parameters that do not directly depend on each other can be updated simultaneously

Parameters that do not directly depend on each other can be updated simultaneously

More precisely, parameters in a **mutually conditionally-independent** set  $W \subseteq S$  can be updated simultaneously. That is,  $W$  satisfying

$$\text{all } w_1, w_2 \in W (w_1 \neq w_2) \text{ satisfy } w_1 \perp\!\!\!\perp w_2 \mid V \setminus W$$

If not all parameters can be collated into a single  $W$ , form a series of  $W$ s and sample in turn.


Define the **topological depth** of a node  $v \in V$  recursively, starting from the nodes with no parents.

$$d(v) = \begin{cases} 0 & \text{if } \text{pa}(v) = \emptyset \\ 1 + \max_{u \in \text{pa}(v)} d(u) & \text{otherwise} \end{cases}$$

# Trivial illustrative example Identifying mutually conditionally-independent sets $W$

Define the **topological depth** of a node  $v \in V$  recursively, starting from the nodes with no parents.


$$d(v) = \begin{cases} 0 & \text{if } \text{pa}(v) = \emptyset \\ 1 + \max_{u \in \text{pa}(v)} d(u) & \text{otherwise} \end{cases}$$



## Trivial illustrative example Identifying mutually conditionally-independent sets $W$

Define the **topological depth** of a node  $v \in V$  recursively, starting from the nodes with no parents.

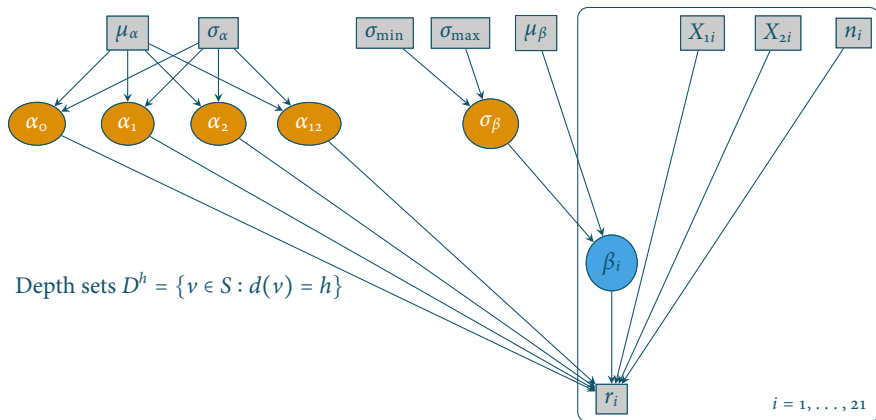
$$d(v) = \begin{cases} 0 & \text{if } \text{pa}(v) = \emptyset \\ 1 + \max_{u \in \text{pa}(v)} d(u) & \text{otherwise} \end{cases}$$



## Trivial illustrative example Identifying mutually conditionally-independent sets $W$

Define the **topological depth** of a node  $v \in V$  recursively, starting from the nodes with no parents.

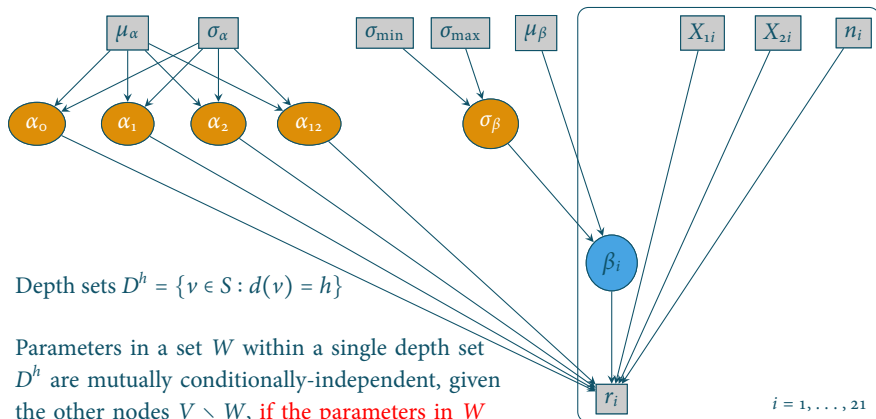
$$d(v) = \begin{cases} 0 & \text{if } \text{pa}(v) = \emptyset \\ 1 + \max_{u \in \text{pa}(v)} d(u) & \text{otherwise} \end{cases}$$



## Trivial illustrative example Identifying mutually conditionally-independent sets $W$

Define the **topological depth** of a node  $v \in V$  recursively, starting from the nodes with no parents.

$$d(v) = \begin{cases} 0 & \text{if } \text{pa}(v) = \emptyset \\ 1 + \max_{u \in \text{pa}(v)} d(u) & \text{otherwise} \end{cases}$$



MultiBUGS aims to

- Parallelise the evaluation of the “likelihood” of ‘fixed effect’-like parameters
- Parallelise sampling of ‘random effect’-like parameters



MultiBUGS aims to

- Parallelise the evaluation of the “likelihood” of ‘fixed effect’-like parameters
- Parallelise sampling of ‘random effect’-like parameters

Let  $\overline{ch} = \text{mean}_{v \in S} |ch(v)|$  be the mean number of children across parameters

### Heuristic algorithm:

Consider each depth set  $D^h$  in turn, starting with the ‘deepest’ set

**if** a parameter has more than  $2 \times \overline{ch}$  children **then**

Parallelise evaluation of this parameter’s “likelihood”

**else**

Sample this parameter in parallel, if possible

**end if**

MultiBUGS aims to

- Parallelise the evaluation of the “likelihood” of ‘fixed effect’-like parameters
- Parallelise sampling of ‘random effect’-like parameters

Let  $\overline{ch} = \text{mean}_{v \in S} |ch(v)|$  be the mean number of children across parameters

### Heuristic algorithm:

Consider each depth set  $D^h$  in turn, starting with the ‘deepest’ set

**if** a parameter has more than  $2 \times \overline{ch}$  children **then**

Parallelise evaluation of this parameter’s “likelihood”

**else**

Sample this parameter in parallel, if possible

**end if**

Users need only specify how many cores they wish to use.

The heuristic is deterministic – makes reproducing the chain easy.

There are 26 stochastic parameters.

---

There are 26 stochastic parameters.

---

Topological depths:

- $d(\beta_1) = \dots = d(\beta_{21}) = 2$
- $d(\alpha_o) = \dots = d(\alpha_{12}) = d(\sigma_\beta) = 1$

There are 26 stochastic parameters.

Topological depths:

- $d(\beta_1) = \dots = d(\beta_{21}) = 2$
- $d(\alpha_o) = \dots = d(\alpha_{12}) = d(\sigma_\beta) = 1$

1. Parameters  $\beta_1, \dots, \beta_{21}$

- Likelihood evaluation not parallelised these parameter have only 1 child and  $\overline{\text{ch}} \approx 4.8$ .
- But,  $\beta_1, \dots, \beta_{21}$  are mutually conditionally-independent and so sampling can be parallelised
- $21 \bmod 4 \neq 0$  so we will have idle cores

Row	Core			
	1	2	3	4
1	$\beta_1$	$\beta_2$	$\beta_3$	$\beta_4$
2	$\beta_5$	$\beta_6$	$\beta_7$	$\beta_8$
3	$\beta_9$	$\beta_{10}$	$\beta_{11}$	$\beta_{12}$
4	$\beta_{13}$	$\beta_{14}$	$\beta_{15}$	$\beta_{16}$
5	$\beta_{17}$	$\beta_{18}$	$\beta_{19}$	$\beta_{20}$
6	$\beta_{21}$			

There are 26 stochastic parameters.

Topological depths:

- $d(\beta_1) = \dots = d(\beta_{21}) = 2$
- $d(\alpha_o) = \dots = d(\alpha_{12}) = d(\sigma_\beta) = 1$

1. Parameters  $\beta_1, \dots, \beta_{21}$

- Likelihood evaluation not parallelised these parameter have only 1 child and  $\overline{\text{ch}} \approx 4.8$ .
- But,  $\beta_1, \dots, \beta_{21}$  are mutually conditionally-independent and so sampling can be parallelised
- $21 \bmod 4 \neq 0$  so we will have idle cores

2. Parameters  $\alpha_o, \alpha_1, \alpha_2, \alpha_{12}$  and  $\sigma_\beta$

- All of these parameters have 21 children  $\overline{\text{ch}} \approx 4.8$ , so likelihood computation is parallelised

Row	Core			
	1	2	3	4
1	$\beta_1$	$\beta_2$	$\beta_3$	$\beta_4$
2	$\beta_5$	$\beta_6$	$\beta_7$	$\beta_8$
3	$\beta_9$	$\beta_{10}$	$\beta_{11}$	$\beta_{12}$
4	$\beta_{13}$	$\beta_{14}$	$\beta_{15}$	$\beta_{16}$
5	$\beta_{17}$	$\beta_{18}$	$\beta_{19}$	$\beta_{20}$
6	$\beta_{21}$			
7	$\alpha_{12}$	$\alpha_{12}$	$\alpha_{12}$	$\alpha_{12}$
8	$\alpha_1$	$\alpha_1$	$\alpha_1$	$\alpha_1$
9	$\alpha_2$	$\alpha_2$	$\alpha_2$	$\alpha_2$
10	$\alpha_o$	$\alpha_o$	$\alpha_o$	$\alpha_o$
11	$\sigma_\beta$	$\sigma_\beta$	$\sigma_\beta$	$\sigma_\beta$

Each core keeps

- The complete DAG, the computation schedule, and associated sampling algorithms
- A copy of the current state of the MCMC
- Two pseudo-random number generation (PRNG) streams
  1. “Core-specific” stream, initialised with a different seed for each core
    - Used when we wish to sample independently across cores
  2. “Common” stream, initialised using the same seed on all cores.
    - Used when we wish to obtain the same samples across cores

Each core keeps

- The complete DAG, the computation schedule, and associated sampling algorithms
- A copy of the current state of the MCMC
- Two pseudo-random number generation (PRNG) streams
  1. “Core-specific” stream, initialised with a different seed for each core
    - Used when we wish to sample independently across cores
  2. “Common” stream, initialised using the same seed on all cores.
    - Used when we wish to obtain the same samples across cores

Most existing code can be reused!



## Implementation notes

Each core keeps

- The complete DAG, the computation schedule, and associated sampling algorithms
- A copy of the current state of the MCMC
- Two pseudo-random number generation (PRNG) streams
  1. “Core-specific” stream, initialised with a different seed for each core
    - Used when we wish to sample independently across cores
  2. “Common” stream, initialised using the same seed on all cores.
    - Used when we wish to obtain the same samples across cores

Most existing code can be reused!

**Likelihood parallelisation:** Just delete the children whose likelihood contribution is calculated elsewhere. MPI function `Allreduce` used to aggregate.

For Metropolis-Hastings: the prior, the sampling of new value, and Metropolis test (redundantly) replicated on every core, using “common” PRNG stream

## Implementation notes

Each core keeps

- The complete DAG, the computation schedule, and associated sampling algorithms
- A copy of the current state of the MCMC
- Two pseudo-random number generation (PRNG) streams
  1. “Core-specific” stream, initialised with a different seed for each core
    - Used when we wish to sample independently across cores
  2. “Common” stream, initialised using the same seed on all cores.
    - Used when we wish to obtain the same samples across cores

Most existing code can be reused!

**Likelihood parallelisation:** Just delete the children whose likelihood contribution is calculated elsewhere. MPI function `Allreduce` used to aggregate.

For Metropolis-Hastings: the prior, the sampling of new value, and Metropolis test (redundantly) replicated on every core, using “common” PRNG stream

**Sampling parallelisation:** Just delete the nodes that are sampled elsewhere from the list of nodes to be updated. Sample using “core-specific” PRNG stream, and propagate new values across cores using `Allgather`.

## Hierarchical regression example The “methadone” model

Based on an analysis linked database of methadone prescriptions given to opioid dependent patients in Scotland (Gao *et al.*, 2016)

425,112 observations, with the following structure:

- Geographic **regions** ( $i = 1, \dots, 8$ )
  - Containing **patients** (20410 in total)
    - Each of whom may have multiple **prescriptions**

For some measurements **patient-level identifiers** are available:

$$y_{ijk} = \sum_{m=1}^4 \beta_m \times \underbrace{x_{mij}}_{\text{covariates}} + \underbrace{u_i}_{\text{region-specific intercept}} + \underbrace{v_i}_{\text{region-specific slope}} \times \underbrace{r_{ijk}}_{\text{covariate}} + \underbrace{w_{ij}}_{\text{patient-level intercept}} + \varepsilon_{ijk}$$

For other measurements **no patient-level identifier** is available:


$$z_{il} = \lambda + \underbrace{u_i}_{\text{region-specific intercept}} + \underbrace{v_i}_{\text{region-specific slope}} \times \underbrace{s_{il}}_{\text{covariate}} + \eta_{il}$$

In OpenBUGS, running chains 2 for 15,000 iterations takes about 32 hours.

JAGS 4.0.0 took 9 hours.

In MultiBUGS 1.0:

- Sampling of pairs of random-effect means and variances parallelised;
- Sampling of person-level random effects  $w_{ij}$  parallelised, except for the component corresponding to the person with the most observations (176 observations)
- The likelihood computation of all the other parameters in the model is parallelised



MultiBUGS 1.0 is available – <https://www.multibugs.org>

MultiBUGS 2.0<sup>1</sup> uses a more efficient system for communicating partial models to cores

Released version requires Windows, but MultiBUGS 2.0 will run on Linux too

Other software:

- JAGS - Martyn Plummer is adopting a similar ideas (using OpenMP)
- NIMBLE? Not sure
- Stan - map/reduce with user-selected sharding

---

<sup>1</sup><https://github.com/MultiBUGS/MultiBUGS>